

ARCHER Software Carpentry bootcamp

Mike Jackson, Arno Proeme, ARCHER CSE Team
michaelj@epcc.ed.ac.uk
aproeme@epcc.ed.ac.uk



software carpentry

| epcc |



Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

http://creativecommons.org/licenses/by-nc-sa/4.0/deed.en_US

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.





EPSRC

NERC
SCIENCE OF THE
ENVIRONMENT

CRAY
THE SUPERCOMPUTER COMPANY

| epcc |



<http://www.archer.ac.uk>
support@archer.ac.uk



software carpentry | epcc |



ARCHER in a nutshell

- UK National Supercomputing Service
 - Replacement for HECToR (they overlap by ~4 months)
- Cray XC30 Hardware
 - Nodes based on 2×Intel Ivy Bridge 12-core processors
 - 64GB (or 128GB) memory per node
 - 3008 nodes in total (72162 cores)
 - Linked by Cray Aries interconnect (dragonfly topology)
- Cray Application Development Environment
 - Cray, Intel, GNU Compilers
 - Cray Parallel Libraries (MPI, SHMEM, PGAS)
 - DDT Debugger, Cray Performance Analysis Tools



software carpentry

epcc



Performance

- HECToR #50 in top 500 with 830 TFlop/s
- ARCHER
 - Designed to provide 3-4 times scientific throughput of HECToR
 - #19 in November 2013 top 500 list with 1.65 PFlop/s
 - Fastest (known) computer in the UK
- Extract the best performance possible from the hardware



Optimization

- “More computing sins are committed in the name of efficiency (without necessarily achieving it) than for any other single reason ..
- ...Including blind stupidity” (Wulf)
- "We should forget about small efficiencies, say about 97% of the time:
- ...Premature optimization is the root of all evil.“ (Knuth 1974)
- “performance variability that derives from differences among programmers of the same language ... is on average as large or larger than the variability found among the different languages.” (Prechelt 2000)



Three rules of optimization

- Don't – optimised code is not readable or maintainable code
 - Don't....yet – designing code is not the same as optimising code
 - Profile first – anything else is just a guess!
-
- How do we prevent ourselves introducing bugs when we're optimizing and parallelizing?
 - Isn't our time more valuable than a computer's time?
 - Hardware life < software life < our life



Optimize ourselves


50%-80% of errors
in

15%-20% of modules
(Davis 1995 quoting Endres
1975, Weinberg 1992)

50% of modules are error free
(Boehm and Basili 2001)

design errors \gg # coding errors
(Boehm et al. 1975)

Technical debt
1-10-100 rule



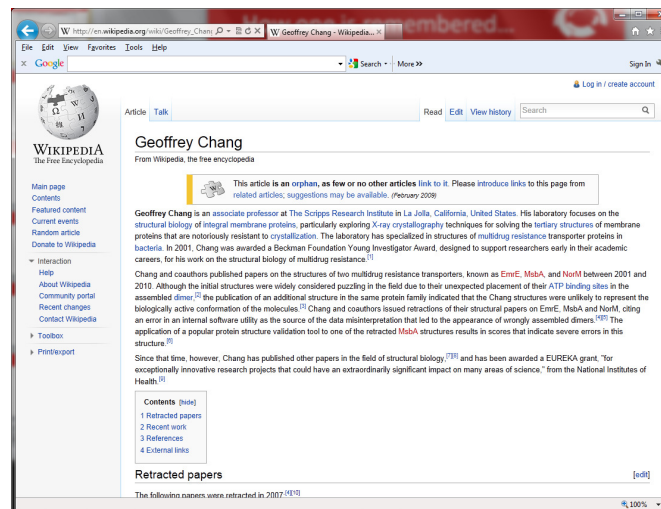
What about correctness?



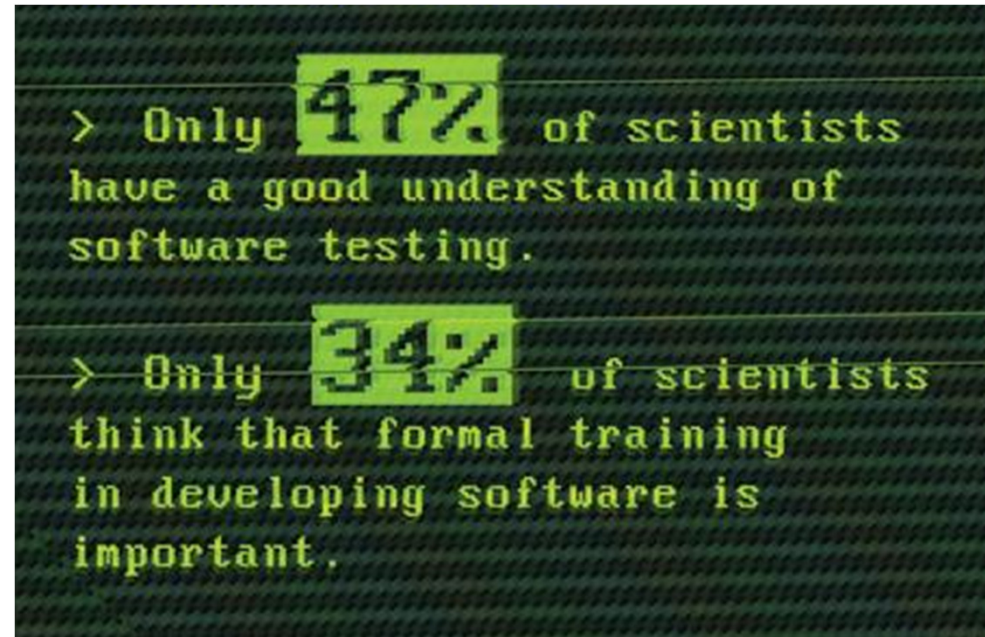

We wish to retract our research article ... and both of our Reports...

An **in-house data reduction program** introduced a **change in sign** for anomalous differences...

Unfortunately, the use of the multicopy refinement procedure still allowed us to obtain reasonable refinement values for the **wrong** structures.



A skills gap



Hannay et al, "How Do Scientists Develop and Use Scientific Software?"

SECSE '09 Proceedings of the 2009 ICSE Workshop on Software Engineering for Computational Science and Engineering, IEEE Computer Society, 2009. DOI: 10.1109/SECSE.2009.5069155

Images courtesy of Greg Wilson and Neil Chue Hong



Filling the gap



DiRAC



Important stuff



Images courtesy of (clockwise from top-left)
oneVillage initiative, Bobbymond, Dajanda,
Anton Novojilov



How to ask for help

- Flag down a helper
- Stick up your sticky note
 - Red – I'm stuck, help!
 - Green – I'm fine

